

D2 Threat Intelligence Command Center

An AI-assisted pipeline that ingests multi-source threat alerts, correlates them into incidents, maps attacker behavior to MITRE ATT&CK, and outputs evidence-grounded BLUF summaries for commanders.

SIEM

SATELLITE

CYBER SENSORS

INTEL REPORTS

THE PROBLEM

Analysts are drowning in disconnected alerts

Defence analysts receive thousands of alerts every day from SIEM systems, satellite feeds, cyber sensors, and intelligence reports — each in a different format. Correlating them manually makes it hard to separate genuine threats from noise, while commanders still need fast, prioritised, BLUF-format assessments.

1000s

of raw alerts a day, across incompatible formats and sources

1

SIEM Systems

Security event logs from network & endpoint tooling

2

Satellite Feeds

Imagery & signal intelligence from orbital assets

3

Cyber Sensors

Intrusion, malware & anomaly detections

4


Intel Reports

Human-authored situation & field reports

OUR SOLUTION

One pipeline: noisy alerts in, a BLUF decision brief out

An AI-powered threat intelligence assistant with a deterministic intelligence engine at its core, and a clearly bounded optional AI-provider interface for future model integration.



BLUF

Bottom Line Up Front — commanders get what happened, why it matters, and what to investigate next, every claim traceable back to source alerts.

KEY FEATURES

What we implemented

1

Multi-Source Ingestion

Normalizes alerts from SIEM, satellite, cyber sensors and intel reports into one common schema.

2

Explainable Correlation

Groups related alerts into incidents and shows the reasoning behind every link it draws.

3

Threat Prioritization

Scores each incident on severity and confidence so the most urgent threats surface first.

4

MITRE ATT&CK Mapping

Matches observed attacker behaviour to ATT&CK techniques, with supporting evidence attached.

5


Evidence-Grounded BLUF

Produces concise Bottom-Line-Up-Front summaries with recommended actions, traceable to source data.

5 / 5

core capabilities shipped in the current prototype

Built for a bounded, extensible AI layer



LANGUAGES




FRAMEWORKS



DATABASE



OTHER



Turning noisy alerts into a trusted, explainable brief

Our strongest work is the evidence-grounded intelligence pipeline: it transforms large volumes of noisy multi-source alerts into prioritised incidents, explains why alerts were correlated, maps attacker behaviour to MITRE ATT&CK, and produces concise BLUF summaries that help commanders understand what happened, why it matters, and what should be investigated.



Explainable

Every correlation and score shows its reasoning — nothing is a black box.



Evidence-Grounded

Every BLUF claim traces back to the source alerts that support it.



Actionable

Outputs prioritised, recommended next steps — not just raw findings.

KNOWN LIMITATIONS

Where this prototype draws the line

Being upfront about scope so judges can weigh the demo correctly.



Synthetic Data

Uses synthetic threat data and a focused MITRE ATT&CK mapping suited for demonstration, not a full production deployment.



No Live IBM Runtime

No IBM Bob or watsonx.ai runtime integration is implemented yet — the intelligence layer exposes a bounded provider interface for a future adapter.



Simplified Security

Operational security controls are simplified for the hackathon environment and would need hardening for real deployment.

THE TEAM

elevateX

Track: AI



Harshil Kalsariya

TEAM LEAD

24ce049@charusat.edu.in



Pranav Dabhi

MEMBER

24ce020@charusat.edu.in



Jenil Viradia

MEMBER

24ce141@charusat.edu.in



Shivam Prajapati

MEMBER

24ce098@charusat.edu.in

Thank You

D2 Threat Intelligence Command Center — elevateX

Source code, setup guide, architecture doc, demo video and screenshots are included in the submission repository under `src/`, `docs/`, `demo/` and `presentation/`.

Questions welcome.

